

wine-quality-prediction

October 13, 2023

1 Predict wine quality using linear regression. Train a model, evaluate, and document insights for accurate predictions in this machine learning task.

2 Author: AHMED REKIK.

2.0.1 THE MACHINE LEARNING INTERNSHIP.

3 MODULES IMPORTATION

```
[69]: import pandas as pd
import numpy as np
import seaborn as sb
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.metrics import mean_squared_error
```

```
[ ]: wine = pd.read_csv("/content/winequalityN.csv")
wine.head()
```

```
[ ]:      type  fixed acidity  volatile acidity  citric acid  residual sugar  \
0  white           7.0           0.27           0.36           20.7
1  white           6.3           0.30           0.34            1.6
2  white           8.1           0.28           0.40            6.9
3  white           7.2           0.23           0.32            8.5
4  white           7.2           0.23           0.32            8.5

      chlorides  free sulfur dioxide  total sulfur dioxide  density  pH  \
0         0.045             45.0             170.0    1.0010  3.00
1         0.049             14.0             132.0    0.9940  3.30
2         0.050             30.0              97.0    0.9951  3.26
3         0.058             47.0             186.0    0.9956  3.19
4         0.058             47.0             186.0    0.9956  3.19

      sulphates  alcohol  quality
0         0.45      8.8       6
1         0.49      9.5       6
```

2	0.44	10.1	6
3	0.40	9.9	6
4	0.40	9.9	6

```
[ ]: wine.describe()
```

```
[ ]:      fixed acidity  volatile acidity  citric acid  residual sugar  \
count      6487.000000      6489.000000  6494.000000      6495.000000
mean         7.216579         0.339691    0.318722         5.444326
std          1.296750         0.164649    0.145265         4.758125
min          3.800000         0.080000    0.000000         0.600000
25%          6.400000         0.230000    0.250000         1.800000
50%          7.000000         0.290000    0.310000         3.000000
75%          7.700000         0.400000    0.390000         8.100000
max         15.900000         1.580000    1.660000        65.800000

      chlorides  free sulfur dioxide  total sulfur dioxide      density  \
count      6495.000000      6497.000000      6497.000000      6497.000000
mean         0.056042      30.525319      115.744574         0.994697
std          0.035036      17.749400       56.521855         0.002999
min          0.009000         1.000000         6.000000         0.987110
25%          0.038000      17.000000       77.000000         0.992340
50%          0.047000      29.000000      118.000000         0.994890
75%          0.065000      41.000000      156.000000         0.996990
max          0.611000      289.000000      440.000000         1.038980

      pH  sulphates  alcohol  quality
count      6488.000000      6493.000000      6497.000000      6497.000000
mean         3.218395         0.531215      10.491801         5.818378
std          0.160748         0.148814         1.192712         0.873255
min          2.720000         0.220000         8.000000         3.000000
25%          3.110000         0.430000         9.500000         5.000000
50%          3.210000         0.510000        10.300000         6.000000
75%          3.320000         0.600000        11.300000         6.000000
max          4.010000         2.000000        14.900000         9.000000
```

```
[6]: wine.shape
```

```
[6]: (6497, 13)
```

```
[10]: wine.isnull().sum()
```

```
[10]: type          0
fixed acidity      10
volatile acidity    8
citric acid         3
residual sugar      2
```

```

chlorides                2
free sulfur dioxide      0
total sulfur dioxide     0
density                  0
pH                       9
sulphates                4
alcohol                  0
quality                  0
dtype: int64

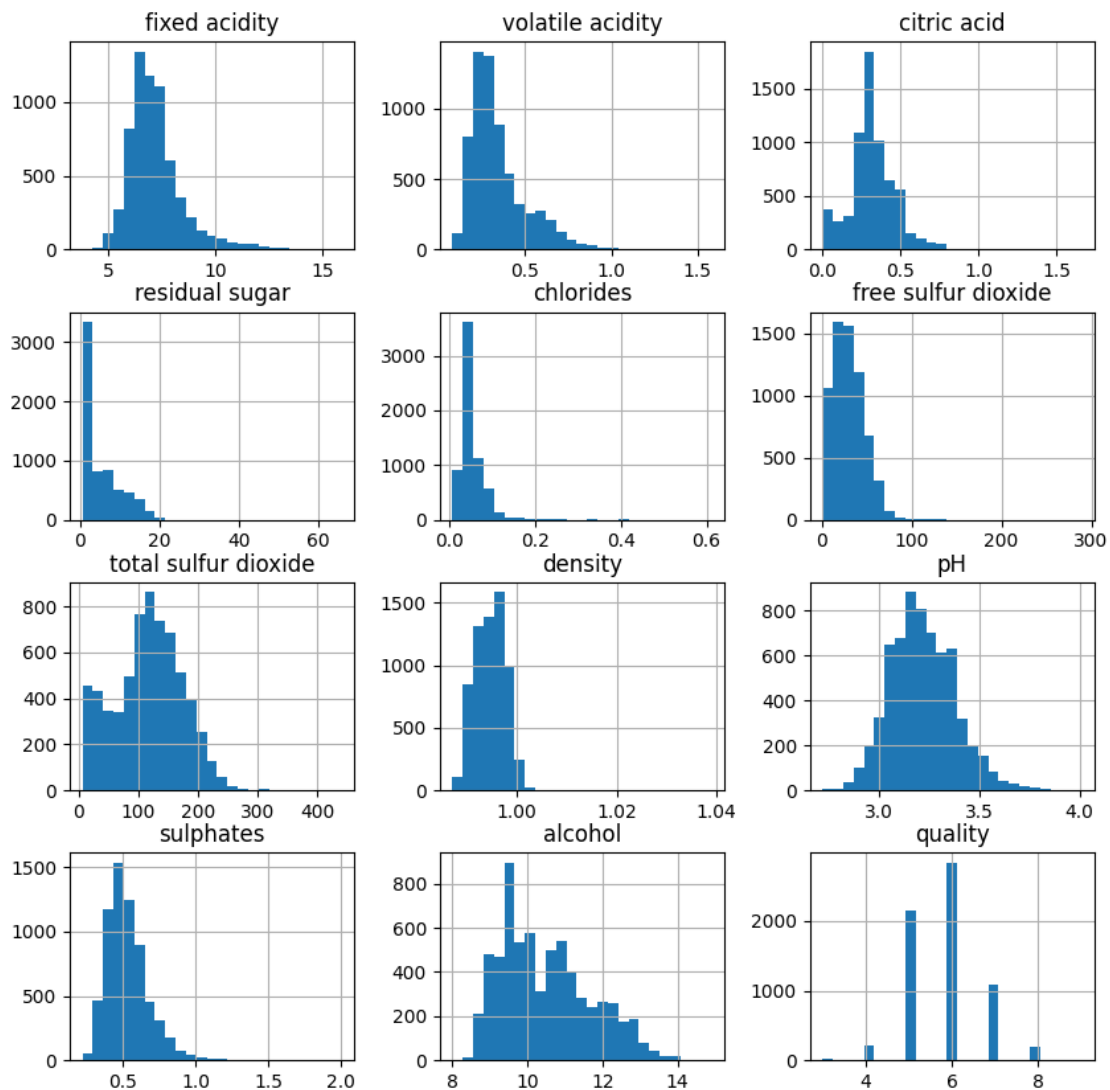
```

4 Visualization

```


[11]: wine.hist(bins=25,figsize=(10,10))
      # display histogram
      plt.show()

```




```
[12]: plt.figure(figsize=[10,6])
      # plot bar graph
      plt.bar(wine['quality'],wine['alcohol'],color='blue')
      # label x-axis
      plt.xlabel('quality')
      #label y-axis
      plt.ylabel('alcohol')
```

```
[12]: Text(0, 0.5, 'alcohol')
```




```
[53]: plt.figure(figsize=(30,15))
      sb.boxplot(data=wine)
```

```
[53]: <Axes: >
```



```
[61]: #Multivariate Analysis
plt.figure(figsize=(30, 15))
sb.pairplot(data=wine, hue='quality')
plt.show()
```

<Figure size 3000x1500 with 0 Axes>




```
[13]: # plotting heatmap
plt.figure(figsize=[19,10],facecolor='blue')
sb.heatmap(wine.corr(),annot=True)
```

<ipython-input-13-b3c2b1f1d9bc>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sb.heatmap(wine.corr(),annot=True)
```

```
[13]: <Axes: >
```



Now, we have to find those features that are fully correlated to each other by this we reduce the number of features from the data.

```
[14]: for a in range(len(wine.corr().columns)):
      for b in range(a):
          if abs(wine.corr().iloc[a,b]) > 0.7:
              name = wine.corr().columns[a]
              print(name)
```

<ipython-input-14-b2e9a4146d97>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
for a in range(len(wine.corr().columns)):
```

<ipython-input-14-b2e9a4146d97>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
if abs(wine.corr().iloc[a,b]) > 0.7:
```

<ipython-input-14-b2e9a4146d97>:4: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
name = wine.corr().columns[a]
```

total sulfur dioxide


```
[16]: new_wine=wine.drop('total sulfur dioxide',axis=1)
```

```
[18]: # plotting heatmap
plt.figure(figsize=[19,10],facecolor='blue')
sb.heatmap(new_wine.corr(),annot=True)
```

<ipython-input-18-6d915c724d26>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sb.heatmap(new_wine.corr(),annot=True)
```

```
[18]: <Axes: >
```



4.1 Handle null values

```
[20]: new_wine.isnull().sum()
```

```
[20]: type          0
fixed acidity     10
volatile acidity   8
citric acid        3
residual sugar     2
chlorides          2
free sulfur dioxide 0
density           0
pH                9
```



```

sulphates          4
alcohol            0
quality            0
dtype: int64

```

```
[22]: new_wine.update(new_wine.fillna(new_wine.mean()))
```

<ipython-input-22-ae10a555ca56>:1: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
new_wine.update(new_wine.fillna(new_wine.mean()))
```

```
[23]: new_wine.isnull().sum()
```

```

[23]: type          0
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
dtype: int64

```

```

[26]: # catogerical vars
next_wine = pd.get_dummies(new_wine,drop_first=True)
# display new dataframe
next_wine

```

```

[26]:      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0              7.0             0.270         0.36             20.7       0.045
1              6.3             0.300         0.34              1.6       0.049
2              8.1             0.280         0.40              6.9       0.050
3              7.2             0.230         0.32              8.5       0.058
4              7.2             0.230         0.32              8.5       0.058
...           ...             ...             ...             ...       ...
6492           6.2             0.600         0.08              2.0       0.090
6493           5.9             0.550         0.10              2.2       0.062
6494           6.3             0.510         0.13              2.3       0.076
6495           5.9             0.645         0.12              2.0       0.075
6496           6.0             0.310         0.47              3.6       0.067

```

	free sulfur dioxide	density	pH	sulphates	alcohol	quality \
0	45.0	1.00100	3.00	0.450000	8.8	6
1	14.0	0.99400	3.30	0.490000	9.5	6
2	30.0	0.99510	3.26	0.440000	10.1	6
3	47.0	0.99560	3.19	0.400000	9.9	6
4	47.0	0.99560	3.19	0.400000	9.9	6
...
6492	32.0	0.99490	3.45	0.580000	10.5	5
6493	39.0	0.99512	3.52	0.531215	11.2	6
6494	29.0	0.99574	3.42	0.750000	11.0	6
6495	32.0	0.99547	3.57	0.710000	10.2	5
6496	18.0	0.99549	3.39	0.660000	11.0	6

	type_white
0	1
1	1
2	1
3	1
4	1
...	...
6492	0
6493	0
6494	0
6495	0
6496	0

[6497 rows x 12 columns]

```
[31]: next_wine["best quality"] = [ 1 if x>=7 else 0 for x in wine.quality]
next_wine
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides \
0	7.0	0.270	0.36	20.7	0.045
1	6.3	0.300	0.34	1.6	0.049
2	8.1	0.280	0.40	6.9	0.050
3	7.2	0.230	0.32	8.5	0.058
4	7.2	0.230	0.32	8.5	0.058
...
6492	6.2	0.600	0.08	2.0	0.090
6493	5.9	0.550	0.10	2.2	0.062
6494	6.3	0.510	0.13	2.3	0.076
6495	5.9	0.645	0.12	2.0	0.075
6496	6.0	0.310	0.47	3.6	0.067

	free sulfur dioxide	density	pH	sulphates	alcohol	quality \
0	45.0	1.00100	3.00	0.450000	8.8	6

1	14.0	0.99400	3.30	0.490000	9.5	6
2	30.0	0.99510	3.26	0.440000	10.1	6
3	47.0	0.99560	3.19	0.400000	9.9	6
4	47.0	0.99560	3.19	0.400000	9.9	6
...
6492	32.0	0.99490	3.45	0.580000	10.5	5
6493	39.0	0.99512	3.52	0.531215	11.2	6
6494	29.0	0.99574	3.42	0.750000	11.0	6
6495	32.0	0.99547	3.57	0.710000	10.2	5
6496	18.0	0.99549	3.39	0.660000	11.0	6

	type_white	best quality
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0
...
6492	0	0
6493	0	0
6494	0	0
6495	0	0
6496	0	0

[6497 rows x 13 columns]

4.1.1 in this dataset ‘Type’ feature contains two types Red and White, where Red consider as 0 and white considers as 1.

5 Splitting dataset

```
[62]: column_train=["fixed acidity","volatile acidity","citric acid","residual_
↪sugar","chlorides","free sulfur_
↪dioxide","density","pH","sulphates","alcohol"]
x=next_wine[column_train]
y=next_wine["quality"]
```

```
[63]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test =train_test_split(x,y,test_size=0.
↪2,random_state=40)
```

6 Normalization

```
[64]: #importing module
from sklearn.preprocessing import MinMaxScaler
# creating normalization object
norm = MinMaxScaler()
# fit data
norm_fit = norm.fit(x_train)
new_xtrain = norm_fit.transform(x_train)
new_xtest = norm_fit.transform(x_test)
# display values
print(new_xtrain)
```

```
[[0.34710744 0.1          0.22289157 ... 0.46511628 0.14044944 0.26086957]
 [0.20661157 0.18          0.15662651 ... 0.30232558 0.15168539 0.11594203]
 [0.24793388 0.06666667 0.18072289 ... 0.21705426 0.16853933 0.14492754]
 ...
 [0.33884298 0.38666667 0.          ... 0.5503876  0.20224719 0.2173913 ]
 [0.24793388 0.34          0.06024096 ... 0.53488372 0.25280899 0.24637681]
 [0.19008264 0.16666667 0.19277108 ... 0.51937984 0.24157303 0.2173913 ]]
```

7 Applying Model

```
[65]: from sklearn.linear_model import LinearRegression
Lrgr= LinearRegression()
Lrgr.fit(x_train,y_train)
```

```
[65]: LinearRegression()
```

```
[67]: # Test prediction (Fixed Acidity, Volatile Acidity, Citric Acid, Residual_
      ↪Sugar, Chlorides, Free Sulfur Dioxide, Total Sulfur Dioxide, Density, pH,
      ↪Sulphates, Alcohol)
print(Lrgr.predict([[15, 0.01, 0, 5, 0.001, 30, 0.95, 3, 0.9, 15]]))
```

```
[9.82662358]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does
not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(
```

8 Evaluation

```
[70]: # Ordinary Least Square (OLS) Principle
x = sm.add_constant(x_train)
model = sm.OLS(y_train, x).fit()
print(model.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          quality    R-squared:                0.285
Model:                  OLS        Adj. R-squared:            0.284
Method:                 Least Squares    F-statistic:            206.7
Date:                  Thu, 12 Oct 2023    Prob (F-statistic):      0.00
Time:                  22:43:11    Log-Likelihood:         -5779.1
No. Observations:      5197    AIC:                    1.158e+04
Df Residuals:          5186    BIC:                    1.165e+04
Df Model:              10
Covariance Type:       nonrobust
=====

```

```

=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----

```

```

const                34.2280      12.837        2.666      0.008      9.063
59.393
fixed acidity         0.0532       0.017        3.096      0.002      0.020
0.087
volatile acidity     -1.3115       0.087       -15.128      0.000     -1.481
-1.142
citric acid          -0.2576       0.086        -2.986      0.003     -0.427
-0.088
residual sugar        0.0315       0.006        5.715      0.000      0.021
0.042
chlorides            -0.2110       0.361        -0.585      0.558     -0.918
0.496
free sulfur dioxide   0.0024       0.001        3.450      0.001      0.001
0.004
density             -33.6139      13.115        -2.563      0.010    -59.324
-7.904
pH                   0.4065       0.100         4.085      0.000      0.211
0.602
sulphates            0.8277       0.085         9.782      0.000      0.662
0.994
alcohol              0.3036       0.018       17.086      0.000      0.269
0.338
=====

```

```

=====
Omnibus:              101.971    Durbin-Watson:           2.014
Prob(Omnibus):        0.000    Jarque-Bera (JB):       216.163
Skew:                 0.009    Prob(JB):               1.15e-47
Kurtosis:             3.999    Cond. No.               6.78e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly

specified.

[2] The condition number is large, $6.78e+04$. This might indicate that there are strong multicollinearity or other numerical problems.

```
[72]: # Evaluate our model using MSE criterion
y_pred = Lrgr.predict(x_test)
print("Mean Squared Error (MSE): ", mean_squared_error(y_test, y_pred))
```

Mean Squared Error (MSE): 0.5694187989872577

```
[73]: # Check quality value from dataset that is listed
print(next_wine['quality'].unique())
```

[6 5 7 8 4 3 9]

```
[74]: # MAPE
def mape(actual, pred):
    actual, pred = np.array(actual), np.array(pred)
    return np.mean(np.abs((actual - pred) / actual)) * 100
```

```
[75]: mape(y_test, y_pred)
# Error of our model is approximately 8-9% (Excellent)
# MAPE result Notes :
# 1. Below 10% = Excellent
# 2. 10% - 20% = Good
# 3. 21% - 50% = Reasonable
# 4. Above 50% = Inaccurate
```

[75]: 10.36550299232918

8.0.1 So the conclusion that we got, our MAPE result is 10.36550299232918 which is approximately 10% (Excellent)

9 THANKS.